

Outline

- Vector space models
- Advantages
- Applications

Why learn vector space models?

Where are you heading?
Where are you from?



Different meaning

What is your age?

How old are you?



Same Meaning

Vector space models applications

- You eat cereal from a bowl
- You buy something and someone else sells it

Vector space models applications

- You eat cereal from a bowl
- You buy something and someone else sells it



Information Extraction



Machine Translation



Chatbots

Fundamental concept

“You shall know a word by the company it keeps”

Firth, 1957



(Firth, J. R. 1957:11)



Summary

- Represent words and documents as **vectors**
- Representation that **captures** relative meaning

Outline

- Co-occurrence —————> Vector representation
- Relationships between words/documents

Word by Word Design

Number of times they *occur together within a certain distance k*

I like simple data

I prefer simple raw data

k=2

	simple	raw	like	I
data	2	1	1	0

Word by Word Design

Number of times they *occur together within a certain distance k*

I like simple data

I prefer simple raw data

k=2

	simple	raw	like	I
data	2	1	1	0
n				

Word by Document Design

Number of times a word *occurs within a certain category*



Corpus

Word by Document Design

Number of times a word *occurs within a certain category*




Word by Document Design

Number of times a word *occurs within a certain category*


	Entertainment	Economy	Machine Learning
Entertainment	500	6620	9320
Economy	7000	4000	1000

Vector Space




	Entertainment	Economy	ML
data	500	6620	9320
film	7000	4000	1000

Vector Space




	Entertainment	Economy	ML
data	500	6620	9320
film	7000	4000	1000

Vector Space



	Entertainment	Economy	ML
data	500	6620	9320
film	7000	4000	1000

Vector Space



	Entertainment	Economy	ML
data	500	6620	9320
film	7000	4000	1000

Measures of “similarity:”
Angle
Distance


Summary

- W/W and W/D, counts of occurrence
- Vector Spaces → Similarity between words/documents

Outline

- Euclidean distance
- N-dimension vector representations comparison


Euclidean distance



Corpus A: (500,7000)

Corpus B: (9320,1000)

Euclidean distance




Corpus A: (500,7000)



Corpus B: (9320,1000)

$$d(B, A) = \sqrt{(B_1 - A_1)^2 + (B_2 - A_2)^2}$$

Euclidean distance




Corpus A: (500,7000)

Corpus B: (9320,1000)

$$d(B, A) = \sqrt{(B_1 - A_1)^2 + (B_2 - A_2)^2}$$

Euclidean distance




Corpus A: (500,7000)

Corpus B: (9320,1000)

$$d(B, A) = \sqrt{(B_1 - A_1)^2 + (B_2 - A_2)^2}$$
$$c^2 = a^2 + b^2$$

$$d(B, A) = \sqrt{(-8820)^2 + (6000)^2}$$

Euclidean distance



Corpus A: (500,7000)

Corpus B: (9320,1000)

$$d(B, A) = \sqrt{(B_1 - A_1)^2 + (B_2 - A_2)^2}$$
$$c^2 = a^2 + b^2$$

$$d(B, A) = \sqrt{(-8820)^2 + (6000)^2}$$

Euclidean distance for n-dimensional vectors

	data	boba	\vec{v} ice-cream
AI	6	0	1
drinks	0	4	6
food	0	6	8

Euclidean distance for n-dimensional vectors

	data	\vec{w}	\vec{v}
AI	6	0	1
drinks	0	4	6
food	0	6	8

Euclidean distance for n-dimensional vectors

	data	\vec{w}	\vec{v}
AI	6	0	1
drinks	0	4	6
food	0	6	8

$$\begin{aligned} &= \sqrt{(1 - 0)^2 + (6 - 4)^2 + (8 - 6)^2} \\ &= \sqrt{1 + 4 + 4} = \sqrt{9} = 3 \end{aligned}$$

Euclidean distance for n-dimensional vectors

	data	\vec{w}	\vec{v}
AI	6	0	1
drinks	0	4	6
food	0	6	8

$$\begin{aligned} &= \sqrt{(1 - 0)^2 + (6 - 4)^2 + (8 - 6)^2} \\ &= \sqrt{1 + 4 + 4} = \sqrt{9} = 3 \end{aligned}$$

$$d(\vec{v}, \vec{w}) = \sqrt{\sum_{i=1}^n (v_i - w_i)^2}$$

Euclidean distance for n-dimensional vectors

	data	\vec{w}	\vec{v}
AI	6	0	1
drinks	0	4	6
food	0	6	8

$$\begin{aligned} &= \sqrt{(1 - 0)^2 + (6 - 4)^2 + (8 - 6)^2} \\ &= \sqrt{1 + 4 + 4} = \sqrt{9} = 3 \end{aligned}$$

$$d(\vec{v}, \vec{w}) = \sqrt{\sum_{i=1}^n (v_i - w_i)^2} \longrightarrow \text{Norm of } (\vec{v} - \vec{w})$$

Euclidean distance in Python

```
# Create numpy vectors v and w
v = np.array([1, 6, 8])
w = np.array([0, 4, 6])

# Calculate the Euclidean distance d
d = np.linalg.norm(v-w)

# Print the result
print("The Euclidean distance between v and w is: ", d)
```

The Euclidean distance between v and w is: 3


Summary

- Straight line between points
- Norm of the difference between vectors


Outline

- Problems with Euclidean Distance
- Cosine similarity


Euclidean distance vs Cosine similarity




Euclidean distance vs Cosine similarity



Euclidean distance vs Cosine similarity



Euclidean distance vs Cosine similarity



Euclidean distance: $d_2 < d_1$

The cosine of the angle
between the vectors

Summary

- Cosine similarity when corpora are different sizes



deeplearning.ai

Cosine Similarity

Outline

- How to get the cosine of the angle between two vectors
- Relation of this metric to similarity

Previous definitions

Vector norm

$$\|\vec{v}\| = \sqrt{\sum_{i=1}^n v_i^2}$$

Previous definitions


Vector norm

$$\|\vec{v}\| = \sqrt{\sum_{i=1}^n v_i^2}$$


Dot product

$$\vec{v} \cdot \vec{w} = \sum_{i=1}^n v_i \cdot w_i$$

Cosine Similarity




Cosine Similarity



$$\hat{v} \cdot \hat{w} = \|\hat{v}\| \|\hat{w}\| \cos(\beta)$$


Cosine Similarity



$$\hat{v} \cdot \hat{w} = \|\hat{v}\| \|\hat{w}\| \cos(\beta)$$

$$\cos(\beta) = \frac{\hat{v} \cdot \hat{w}}{\|\hat{v}\| \|\hat{w}\|}$$

Cosine Similarity





$$\hat{v} \cdot \hat{w} = \|\hat{v}\| \|\hat{w}\| \cos(\beta)$$

$$\cos(\beta) = \frac{\hat{v} \cdot \hat{w}}{\|\hat{v}\| \|\hat{w}\|}$$

$$\begin{aligned} &= \frac{(20 \times 30) + (40 \times 20)}{\sqrt{20^2 + 40^2} \times \sqrt{30^2 + 20^2}} \\ &= 0.87 \end{aligned}$$

Cosine Similarity



Summary

- Cosine \propto Similarity
- Cosine Similarity gives values between 0 and 1

Outline

- How to use vector representations

Manipulating word vectors



USA



Washington
DC




Russia




?

Manipulating word vectors




[Mikolov et al, 2013, Distributed Representations of Words and Phrases and their Compositionality]

Manipulating word vectors




[Mikolov et al, 2013, Distributed Representations of Words and Phrases and their Compositionality]

Manipulating word vectors



[Mikolov et al, 2013, Distributed Representations of Words and Phrases and their Compositionality]

Manipulating word vectors



$$\text{Washington} - \text{USA} = [5 \quad -1]$$

$$\text{Russia} + [5 \quad -1] = [10 \quad 4]$$



Moscow

[Mikolov et al, 2013, Distributed Representations of Words and Phrases and their Compositionality]

Summary

- Use known relationships to make predictions

Outline

- Some motivation for visualization

Outline

- Some motivation for visualization
- Principal Component Analysis

Visualization of word vectors

oil	0.20	...	0.10
gas	2.10	...	3.40
city	9.30	...	52.1
town	6.20	...	34.3

Visualization of word vectors

$d > 2$

oil	0.20	...	0.10
gas	2.10	...	3.40
city	9.30	...	52.1
town	6.20	...	34.3



oil & gas



town & city

Visualization of word vectors

$d > 2$

oil	0.20	...	0.10
gas	2.10	...	3.40
city	9.30	...	52.1
town	6.20	...	34.3



oil & gas



town & city

How can you visualize if your representation captures these relationships?

Visualization of word vectors


$d > 2$

oil	0.20	...	0.10
gas	2.10	...	3.40
city	9.30	...	52.1
town	6.20	...	34.3


Visualization of word vectors

$d > 2$				$d = 2$			
oil	0.20	...	0.10	oil	2.30	21.2	
gas	2.10	...	3.40	gas	1.56	19.3	
city	9.30	...	52.1	city	13.4	34.1	
town	6.20	...	34.3	town	15.6	29.8	


Visualization of word vectors



Visualization of word vectors



Principal Component Analysis





Summary

- Original Space  Uncorrelated features  Dimension reduction
- Visualization to see words relationships in the vector space

Outline

- How to get uncorrelated features
- How to reduce dimensions while retaining as much information as possible

Principal Component Analysis




PCA algorithm

Eigenvector: Uncorrelated features for your data

Eigenvalue: the amount of information retained by each feature

PCA algorithm




Mean Normalize Data $x_i = \frac{x_i - \mu_{x_i}}{\sigma_{x_i}}$

Get Covariance Matrix Σ

Perform SVD $SVD(\Sigma)$

PCA algorithm



Mean Normalize Data $x_i = \frac{x_i - \mu_{x_i}}{\sigma_{x_i}}$


Get Covariance Matrix Σ

Perform SVD $SVD(\Sigma)$


Eigenvalues



Eigenvectors




PCA algorithm



Eigenvectors	Eigenvalues
$\begin{pmatrix} & & \\ U & \dots & \end{pmatrix}$	$\begin{pmatrix} \cdot & \cdot & \cdot \\ \cdot & S & \cdot \\ \cdot & \cdot & \cdot \end{pmatrix}$

PCA algorithm



Eigenvectors Eigenvalues

$$\begin{pmatrix} | & | & | \\ U & \dots & | \end{pmatrix} \quad \begin{pmatrix} \cdot & \cdot & \cdot \\ S & \dots & \cdot \end{pmatrix}$$

Dot Product to
Project Data

$$X' = XU[:, 0 : 2]$$

Percentage of
Retained Variance

$$\frac{\sum_{i=0}^1 S_{ii}}{\sum_{j=0}^d S_{jj}}$$

Summary

- Eigenvectors give the direction of uncorrelated features
- Eigenvalues are the variance of the new features
- Dot product gives the projection on uncorrelated features